

---

# **Pod Network Bridge Audit Report**

---

Prepared by Riley Holterhus

January 29, 2026

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	About Pod Network . . . . .	3
1.2	About the Auditor . . . . .	3
1.3	Disclaimer . . . . .	3
<b>2</b>	<b>Audit Overview</b>	<b>4</b>
2.1	Scope of Work . . . . .	4
2.2	Summary of Findings . . . . .	4
<b>3</b>	<b>Findings</b>	<b>5</b>
3.1	Medium Severity Findings . . . . .	5
3.1.1	Version updates require careful merkle tree construction . . . . .	5
3.2	Low Severity Findings . . . . .	6
3.2.1	Upgradeable contract inherits base AccessControl . . . . .	6
3.2.2	minAmount and claimLimit checks can result in unclaimable funds . . . . .	6
3.3	Informational Findings . . . . .	6
3.3.1	Invariant between validatorCount and adversarialResilience is split across functions . . . . .	6
3.3.2	Limit behavior is affected by order splitting . . . . .	7
3.3.3	lenData can be used instead of a hardcoded length . . . . .	7
3.3.4	Usage resets favor faster claimers . . . . .	8

# **1 Introduction**

## **1.1 About Pod Network**

Pod Network is a Layer-1 network with an architecture that differs from most traditional blockchains. One notable aspect is that transactions are partially ordered rather than enforced under a strict total order, which allows for faster transaction processing.

## **1.2 About the Auditor**

Riley Holterhus is an independent security researcher who focuses on Solidity smart contracts. Other than conducting independent security reviews, he works as a Lead Security Researcher at [Spearbit](#), and also searches for vulnerabilities in live codebases. Riley can be reached by email at [rileyholterhus@gmail.com](mailto:rileyholterhus@gmail.com), by Telegram at [@holterhus](#) and on Twitter/X at [@rileyholterhus](#).

## **1.3 Disclaimer**

This report is intended to detail the identified vulnerabilities of the reviewed smart contracts and should not be construed as an endorsement or recommendation of any project, individual or entity. While the authors have made reasonable efforts to detect potential issues, the absence of any undetected vulnerabilities or issues cannot be guaranteed. Additionally, the security of the smart contracts may be affected by future changes or updates. By using the information in this report, you acknowledge that you are doing so at your own risk and that you should exercise your own judgment when implementing any recommendations or making decisions based on the findings. This report has been provided on an “as-is” basis and DOES NOT CONSTITUTE A GUARANTEE OR WARRANTY OF ANY FORM.

## 2 Audit Overview

### 2.1 Scope of Work

From January 24, 2026 to January 25, 2026, Riley Holterhus conducted a manual security review of Pod's Bridge smart contract. The goal was to identify potential vulnerabilities and logic issues in the codebase.

The audit was performed on the codebase found in the [podnetwork/pod-sdk](#) GitHub repository. The audit started on commit [078fbab](#). The files in scope for the audit were:

- protocol/src/Bridge.sol
- protocol/src/lib/ProofLib.sol

The Bridge contract is intended to be deployed on an EVM chain such as Ethereum mainnet or Base, and is responsible for bridging tokens between the EVM chain and Pod. Users deposit ERC20 tokens into the contract, which is observed by a validator set and results in corresponding tokens being minted on Pod. When bridging tokens back from Pod, users receive a set of signatures (or an equivalent proof against an admin-provided merkle tree) that can be submitted to the contract to transfer tokens out. The contract also includes migration logic to allow the admin to move funds into a new contract in the event of an issue.

### 2.2 Summary of Findings

Each finding from the audit has been assigned a severity level of "Critical", "High", "Medium", "Low" or "Informational". These severities aim to capture the impact and likelihood of each potential issue.

In total, **7 findings** were identified: 1 Medium, 2 Low, and 4 Informational. The description and status of the findings are provided in the sections below.

## 3 Findings

### 3.1 Medium Severity Findings

#### 3.1.1 Version updates require careful merkle tree construction

**Description:** The Bridge allows the admin to update the version number, which updates the domainSeparator and causes all subsequent txHash calculations to be derived using the new version. There is a comment above updateValidatorConfig() that references this behavior:

```
/**  
 * ...  
 * @dev Admin may keep the same version if the config update doesn't invalidate past certificates.  
 *       When the version is updated, past certificates become invalid and only merkle proofs work  
 *       for claims from before the version change.  
 * ...  
 */
```

While this comment indicates that merkle proofs are intended to be the mechanism for proving unclaimed transactions from a previous version, the exact mechanics of how this is expected to work are more nuanced.

In order for merkle proofs to be used for transactions from a previous version, an updated merkle root would need to be submitted that includes all previously unclaimed transactions, with their leaves recomputed using the new domain separator. Since this effectively reintroduces the same transactions under new txHash values, the update process must be careful not to include any transactions that were already claimed under the old version, as they would become claimable again under the new txHash in the merkle tree.

So, for example, it's important that the version upgrade and the updating of the merkle root are not performed in a single step based on an assumed set of unclaimed transactions. If both were updated simultaneously, an attacker could front-run the update by claiming a transaction under the old version, and then claim the same transaction again after the update using the new merkle root. To avoid this race condition, the version update and merkle root update would need to be performed as separate steps. This approach could be error-prone and is currently not documented in the code.

**Recommendation:** Consider changing this behavior to be more robust. One possible approach would be to allow the claimer to specify the version associated with the transaction being claimed. This would allow the txHash for a given transaction to remain the same across version updates, which in turn would allow merkle root updates to include transactions that were already claimed, without introducing a risk of double claims.

**Pod:** Fixed in [commit 6e7293b](#).

**Auditor:** Verified.

## 3.2 Low Severity Findings

### 3.2.1 Upgradeable contract inherits base AccessControl

**Description:** The Bridge contract is intended to be upgradeable behind a proxy, but it is inheriting OpenZeppelin's base AccessControl contract rather than AccessControlUpgradeable from the upgradeable OpenZeppelin codebase.

In practice, using AccessControl behind a proxy doesn't introduce any immediate issues, as it doesn't have any special constructor or initialization logic. However using AccessControlUpgradeable would be more consistent and also more forward-compatible, since the upgradeable version uses an ERC-7201 storage layout.

**Recommendation:** Consider inheriting the AccessControlUpgradeable contract instead of AccessControl.

**Pod:** Fixed in commit [8aeedab](#).

**Auditor:** Verified. The AccessControlUpgradeable contract is now being inherited.

### 3.2.2 minAmount and claimLimit checks can result in unclaimable funds

**Description:** Every token in the Bridge is configured with minAmount, depositLimit, and claimLimit thresholds.

For claims, these thresholds can theoretically create unclaimable situations. For example, if a user burns tokens on Pod to bridge back to mainnet with an amount below minAmount, or if minAmount is increased after the burn, the user wouldn't be able to claim their tokens through the Bridge. Another example is if a user bridges back an amount larger than the token's claimLimit. This is unclaimable because claims can't be split up into smaller pieces once they're initiated and signed.

**Recommendation:** Ensure users do not initiate bridge operations back to mainnet with amounts below minAmount or above the claimLimit, as this would result in stuck funds. If possible, enforce these thresholds at the time the bridge is initiated on the Pod side. Additionally, ensure that later increases to minAmount or decreases to claimLimit do not create stuck funds.

**Pod:** Acknowledged.

**Auditor:** Acknowledged.

## 3.3 Informational Findings

### 3.3.1 Invariant between validatorCount and adversarialResilience is split across functions

**Description:** The \_addRemoveValidators() and \_setAdversarialResilience() functions are admin helper functions. The \_addRemoveValidators() function adds and removes validators from the active validator set, while \_setAdversarialResilience() updates the adversarialResilience value, which must always be less than or equal to validatorCount.

Currently, `_addRemoveValidators()` does not verify that changes to the validator set preserve the invariant `validatorCount >= adversarialResilience`. Instead, this invariant is enforced indirectly by relying on a subsequent call to `_setAdversarialResilience()`. As a result, future changes to the code must ensure that `_addRemoveValidators()` is always followed by `_setAdversarialResilience()` to avoid violating this invariant.

**Recommendation:** Consider documenting this behavior for future versions of the code. Alternatively, consider combining `_addRemoveValidators()` and `_setAdversarialResilience()` into a single function so that the invariant is upheld more explicitly.

**Pod:** Fixed in [commit a64ef9c](#).

**Auditor:** Verified.

### 3.3.2 Limit behavior is affected by order splitting

**Description:** The Bridge contract enforces both deposit and claim limits using a rolling usage tracker. Before each deposit or claim, the current usage is incremented by the requested amount and checked against the limit, as shown below:

```
uint256 newConsumed = usage.consumed + amount;
if (newConsumed > maxTotalAmount) {
    if (block.timestamp < usage.lastUpdated + 1 days || amount > maxTotalAmount) {
        revert DailyLimitExhausted();
    } else {
        usage.lastUpdated = block.timestamp;
        usage.consumed = amount;
    }
} else {
    usage.consumed = newConsumed;
}
```

Notice that if an action pushes the current usage past the limit, the usage is reset and the full amount is applied to the new period, causing any remaining capacity from the previous period to be lost. As a result, splitting an action into multiple smaller actions near a reset can lower the end usage, even when the overall amount is the same.

**Recommendation:** Consider changing the behavior so that splitting an action does not lead to a different end result. This can be achieved by filling remaining capacity first and only applying excess amounts to the usage in a new period.

**Pod:** Fixed in [commit e864c76](#).

**Auditor:** Verified. The code now fills remaining capacity first and applies the excess to the new period.

### 3.3.3 `lenData` can be used instead of a hardcoded length

**Description:** In the `_depositTxHash()` function, `dataHash` is computed using a hardcoded length:

```
dataHash := keccak256(ptr, 0x64)
```

Note that this hardcoded `0x64` value already corresponds to the `lenData` variable defined in the function.

**Recommendation:** Consider replacing the hardcoded `0x64` value with the `lenData` variable.

**Pod:** Fixed in [commit 9fa3752](#).

**Auditor:** Verified.

### 3.3.4 Usage resets favor faster claimers

**Description:** The Bridge enforces usage limits on both deposits and claims, which can be reset every 24 hours. After a reset, usage is allocated on a first-come first-served basis. This means that when usage resets, multiple users might rush to consume the new capacity.

This behavior could negatively affect the user-experience, especially for claims. For claims, bridging back to mainnet is a two-step process where users first burn their tokens on Pod and then submit their claim to the Bridge later. If the claim capacity is consistently consumed immediately after a reset, slower users may have to wait multiple days to claim, despite already having burned their tokens.

**Recommendation:** Keep this behavior in mind, and if it causes issues in practice, consider implementing a queue mechanism.

**Pod:** Acknowledged.

**Auditor:** Acknowledged.